

INTRODUCTION TO STAN

BAYESIAN STATISTICS FOR ECOLOGISTS

IGB 18. TO 26. NOVEMBER 2019

WHAT IS STAN?

- ▶ General purpose modelling and scientific programming language
- ▶ Imperative (like R, unlike BUGS) and probabilistic language
- ▶ Fully Bayesian inference using Hamiltonian Monte Carlo
 - ▶ Your model is compiled to a C++ program that is executed from within R
 - ▶ HMC converges faster and is more stable than Gibbs in many cases
 - ▶ Lower autocorrelation than Gibbs

WHY STAN?

- ▶ In many cases faster than Gibbs, even Gibbs in C++
- ▶ Much of the tuning, warmup, care and feeding is done automatically
- ▶ Model specification is very concise

VARIABLES

- ▶ Stan is **strongly typed**: all variables must be declared and used in their proper scope

- ▶ real, int for single numbers
- ▶ vector[i], matrix[i, j] for real-valued arrays to be used with linear algebra
- ▶ arrays (for all other sequences)

- ▶ Stan supports variable constraints

```
real my_real;  
int my_int;
```

```
vector [10] my_vector;  
matrix [10, 20] my_matrix;
```

```
int my_int_array [10];  
vector [10] array_of_vectors [5, 2];
```

```
real <lower=0, upper=1> constrained_real;
```

STRUCTURE OF A PROGRAM

- ▶ A Stan program consists of a series of optional **blocks** delimited by curly braces
 - `functions {}`
 - `data {}`
 - `transformed data {}`
 - `parameters {}`
 - `transformed parameters {}`
 - `model {}`
 - `generated quantities {}`
- ▶ Order matters, and variables are visible from the top down (**model** can see variables declared in **data**, but not vice-versa!)
- ▶ Variable declarations must come first in each block
- ▶ **data**: only declares the data variables that will be passed to Stan
- ▶ **parameters**: only declares parameters that will be used in the model
- ▶ **model**: defines the log probability of the model

STAN PRACTICE—TREE MORTALITY

- ▶ Often helpful to work backwards, start with the model definition
- ▶ `died` is a data vector giving the number of trees that died
- ▶ `ntrees` is the vector of the total number of trees
- ▶ `theta` is the parameter to estimate
- ▶ To do:
 - ▶ Fill in the missing blocks
 - ▶ Add prior
 - ▶ Run the model in R

STAN PRACTICE—TREE MORTALITY

- ▶ Often helpful to work backwards, start with the model definition
- ▶ `died` is a data vector giving the number of trees that died
- ▶ `ntrees` is the vector of the total number of trees
- ▶ `theta` is the parameter to estimate
- ▶ To do:
 - ▶ Fill in the missing blocks
 - ▶ Add prior
 - ▶ Run the model in R

```
data {  
    // fill in data declarations  
}  
parameters {  
    // fill in parameters  
}  
model {  
    died ~ binomial(ntrees, theta);  
}
```

STAN PRACTICE—TREE MORTALITY

- ▶ It is helpful to include the number of data points as a data variable
- ▶ All data and parameters must be declared
- ▶ **Data constraints** serve as error checking, in this case checking for negative numbers
- ▶ **Parameter constraints** prevent the sampler from selecting impossible values
- ▶ All **unobserved** objects (e.g., theta) should appear on the left side of a statement:
either = or ~
- ▶ Check carefully! It's not an error to forget this step

STAN PRACTICE—TREE MORTALITY

- ▶ It is helpful to include the number of data points as a data variable
- ▶ All data and parameters must be declared
- ▶ **Data constraints** serve as error checking, in this case checking for negative numbers
- ▶ **Parameter constraints** prevent the sampler from selecting impossible values
- ▶ All **unobserved** objects (e.g., θ) should appear on the left side of a statement:
either $=$ or \sim
- ▶ Check carefully! It's not an error to forget this step

```
data {  
    int<lower=0> n; // number of data points  
    int<lower=0> died[n];  
    int<lower=0> ntrees[n];  
}  
parameters {  
    real<lower=0, upper=1> theta;  
}  
model {  
    died ~ binomial(ntrees, theta);  
    theta ~ ?????  
}
```

STAN PRACTICE—TREE MORTALITY

- ▶ Running the model is simply a matter of pointing the `stan()` function at the `.stan` file you created
- ▶ Examine the summary of the model. Is the parameter estimate different from the MLE?
- ▶ You can use `bayesplot` to produce some useful diagnostics

