


BAYESIAN GLMS

BAYESIAN STATISTICS FOR ECOLOGISTS

IGB 18. TO 26. NOVEMBER 2019


A BAYESIAN REGRESSION MODEL

- ▶ Goal: from observations of an i.i.d variable y and some covariate x , estimate the slope and intercept of the relationship between x and y



A BAYESIAN REGRESSION MODEL


- ▶ Goal: from observations of an i.i.d variable y and some covariate x , estimate the slope and intercept of the relationship between x and y



A BAYESIAN REGRESSION MODEL

- ▶ Goal: from observations of an i.i.d variable y and some covariate x , estimate the slope and intercept of the relationship between x and y

$$\hat{y} = \alpha + \beta x$$




A BAYESIAN REGRESSION MODEL


- ▶ Goal: from observations of an i.i.d variable y and some covariate x , estimate the slope and intercept of the relationship between x and y

$$\hat{y} = \alpha + \beta x$$

$$y \sim \mathbb{D}(\cdot)$$




A BAYESIAN REGRESSION MODEL



A BAYESIAN REGRESSION MODEL

- ▶ We need to define some reasonable distribution for y


$$\hat{y} = \alpha + \beta x$$



A BAYESIAN REGRESSION MODEL

- ▶ We need to define some reasonable distribution for y

$$\hat{y} = \alpha + \beta x$$
$$y \sim N(\hat{y}, \sigma)$$




A BAYESIAN REGRESSION MODEL

- ▶ We need to define some reasonable distribution for y

$$\hat{y} = \alpha + \beta x$$
$$y \sim N(\hat{y}, \sigma)$$

- ▶ y is a random variable, distributed normally with a mean equal to the regression line



BAYESIAN REGRESSION (R)

- ▶ The likelihood simply translates this idea into code
- ▶ Remember to use logs and to sum across the entire dataset
- ▶ What prior makes sense?

$$\hat{y} = \alpha + \beta x$$

$$y \sim \mathcal{N}(\hat{y}, \sigma)$$

$$pr(y|x, \alpha, \beta, \sigma)$$

```
log_lik <- function(alpha, beta, sigma, x, y) {  
  y_hat <- alpha + beta * x  
  ll <- sum(dnorm(y, y_hat, sigma, log=TRUE))  
  return(ll)  
}
```

```
log_prior <- function(alpha, beta, sigma) {  
  
}
```

BAYESIAN REGRESSION (R)

- ▶ The likelihood simply translates this idea into code
- ▶ Remember to use logs and to sum across the entire dataset
- ▶ What prior makes sense?

```
log_lik <- function(alpha, beta, sigma, x, y) {  
  y_hat <- alpha + beta * x  
  ll <- sum(dnorm(y, y_hat, sigma, log=TRUE))  
  return(ll)  
}
```

```
log_prior <- function(alpha, beta, sigma) {  
  lp <- dnorm(alpha, 0, 50, log = TRUE)  
  lp <- lp + dnorm(beta, 0, 25, log=TRUE)  
  lp <- lp + dexp(sigma, 0.1, log=TRUE)  
  return(lp)  
}
```

BAYESIAN REGRESSION LIKELIHOOD (STAN)

- ▶ The likelihood simply translates this idea into code
- ▶ In Stan, the we have to declare all variables
- ▶ No need to worry about taking the log of the likelihood, Stan takes care of this

```
data {  
    int<lower=0> n; // number of data points  
    vector[n] x;  
    vector[n] y;  
}  
parameters {  
    real<lower=0> sigma;  
    real alpha;  
    real bet;  
}  
transformed parameters {  
    vector[n] y_hat;  
    y_hat = alpha + bet * x;  
}  
model {  
    y ~ normal(y_hat, sigma);  
}
```

BAYESIAN REGRESSION LIKELIHOOD (STAN)

- ▶ The likelihood simply translates this idea into code
- ▶ In Stan, the we have to declare all variables
- ▶ No need to worry about taking the log of the likelihood, Stan takes care of this
- ▶ Prior as before from R

```
transformed parameters {  
    vector[n] y_hat;  
    y_hat = alpha + bet * x;  
}  
model {  
    y ~ normal(y_hat, sigma);  
    alpha ~ normal(0, 50);  
    bet ~ normal(0, 25);  
    sigma ~ exponential(0.1);  
}
```

THE GENERALISED LINEAR MODEL

$$y \sim \mathcal{N}(\hat{y}, \sigma)$$

THE GENERALISED LINEAR MODEL

- ▶ If we have two x-variables, the model becomes

$$y \sim N(\hat{y}, \sigma)$$

THE GENERALISED LINEAR MODEL

$$y \sim N(\hat{y}, \sigma)$$

- ▶ If we have two x-variables, the model becomes

$$\hat{y} = \alpha + \beta_1 x_1 + \beta_2 x_2$$

THE GENERALISED LINEAR MODEL

$$y \sim N(\hat{y}, \sigma)$$

- ▶ If we have two x-variables, the model becomes

$$\hat{y} = \alpha + \beta_1 x_1 + \beta_2 x_2$$

- ▶ Which generalises to a **matrix** of x-variables and a **vector** of β s

THE GENERALISED LINEAR MODEL

$$y \sim N(\hat{y}, \sigma)$$

- ▶ If we have two x-variables, the model becomes

$$\hat{y} = \alpha + \beta_1 x_1 + \beta_2 x_2$$

- ▶ Which generalises to a **matrix** of x-variables and a **vector** of β s

$$\hat{y} = \alpha + \boldsymbol{\beta} \mathbf{x}$$

THE GENERALISED LINEAR MODEL

$$\mathbb{L}(\hat{y}) = \alpha + \beta \mathbf{x}$$

$$y \sim \mathbb{D}(\mu = \hat{y}, \dots)$$

THE GENERALISED LINEAR MODEL

- ▶ Further, there is no requirement that the observed y 's be distributed **normally**, nor that the relationship between \hat{y} and \mathbf{x} be **strictly linear**

$$\mathbb{L}(\hat{y}) = \alpha + \beta \mathbf{x}$$

$$y \sim \mathbb{D}(\mu = \hat{y}, \dots)$$

THE GENERALISED LINEAR MODEL

$$\mathbb{L}(\hat{y}) = \alpha + \beta \mathbf{x}$$

$$y \sim \mathbb{D}(\mu = \hat{y}, \dots)$$

THE GENERALISED LINEAR MODEL


- ▶ The choice of **error distribution** and **link function** will depend on the problem
- ▶ For y's representing counts, a Poisson error distribution (likelihood) combined with a **log-link** can make sense
- ▶ For counts of successes given a number of trials, we might use **Binomial** with the **logit** or **probit** link

$$\mathbb{L}(\hat{y}) = \alpha + \beta \mathbf{x}$$

$$y \sim \mathbb{D}(\mu = \hat{y}, \dots)$$





GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ What are the posterior credible intervals for the parameters?
- ▶ How to interpret this in terms of the effect of temperature on mortality?







GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ Traces seem ok
- ▶ Posteriors look normal



GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ Traces seem ok
- ▶ Posteriors look normal
- ▶ Multivariate normal, with strong correlation ($r = 0.98$)



GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ What are the posterior credible intervals for the parameters?

```
>apply(samples, 2, quantile, c(0.05, 0.95))
```

parameters

intercept


slope

5% 0.7901569 -0.6613466

95% 2.1106432 -0.4294134

GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ How to interpret this in terms of the effect of temperature on mortality?



GLM IN STAN/MCMC

- ▶ How does temperature affect the mortality of *Tsuga canadensis*?
- ▶ How to interpret this in terms of the effect of temperature on mortality?

